

# Class 08: Machine Learning Mini Project

Bianca Barriga

Function focus for today: `grep()`, `kmeans()`, `hclust()`, `prcomp()`

#Import the dataset

Before we can begin our analysis we first have to download and import our data correctly into our R session.

```
wisc.df <- read.csv("WisconsinCancer.csv", row.names=1)
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
842302	0.11840	0.27760	0.3001	0.14710
842517	0.08474	0.07864	0.0869	0.07017
84300903	0.10960	0.15990	0.1974	0.12790
84348301	0.14250	0.28390	0.2414	0.10520
84358402	0.10030	0.13280	0.1980	0.10430
843786	0.12780	0.17000	0.1578	0.08089

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585
84348301	0.2597	0.09744	0.4956	1.1560	3.445
84358402	0.1809	0.05883	0.7572	0.7813	5.438
843786	0.2087	0.07613	0.3345	0.8902	2.217

	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
--	---------	---------------	----------------	--------------	-------------------

842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137
symmetry_se fractal_dimension_se radius_worst texture_worst					
842302	0.03003	0.006193	25.38	17.33	
842517	0.01389	0.003532	24.99	23.41	
84300903	0.02250	0.004571	23.57	25.53	
84348301	0.05963	0.009208	14.91	26.50	
84358402	0.01756	0.005115	22.54	16.67	
843786	0.02165	0.005082	15.47	23.75	
perimeter_worst area_worst smoothness_worst compactness_worst					
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	
84358402	152.20	1575.0	0.1374	0.2050	
843786	103.40	741.6	0.1791	0.5249	
concavity_worst concave.points_worst symmetry_worst					
842302	0.7119	0.2654	0.4601		
842517	0.2416	0.1860	0.2750		
84300903	0.4504	0.2430	0.3613		
84348301	0.6869	0.2575	0.6638		
84358402	0.4000	0.1625	0.2364		
843786	0.5355	0.1741	0.3985		
fractal_dimension_worst					
842302	0.11890				
842517	0.08902				
84300903	0.08758				
84348301	0.17300				
84358402	0.07678				
843786	0.12440				

Q1. How many samples are in this dataset?

```
nrow(wisc.df)
```

```
[1] 569
```

```
ncol(wisc.df)
```

```
[1] 31
```

Q. How many variables (columns)?

```
ncol(wisc.df)
```

```
[1] 31
```

Q2. How many M and B samples are there?

```
table(wisc.df$diagnosis)
```

```
  B    M  
357 212
```

Q3. How many variables/features in the data are suffixed with `_mean`?

```
txt <- c("_mean")  
length(grep(txt, colnames(wisc.df), value =TRUE))
```

```
[1] 10
```

Q. what features are “mean” values?

```
txt <- c("_mean")  
grep(txt, colnames(wisc.df), value =TRUE)
```

```
[1] "radius_mean"          "texture_mean"         "perimeter_mean"  
[4] "area_mean"            "smoothness_mean"      "compactness_mean"  
[7] "concavity_mean"       "concave.points_mean"  "symmetry_mean"  
[10] "fractal_dimension_mean"
```

I need to remove the first diagnosis column from my data before doing any analysis. I will store it for later.

```
# We can use -1 here to remove the first column
wisc.data <- wisc.df[,-1]
```

```
# Create diagnosis vector for later
diagnosis <- as.factor(wisc.df$diagnosis)
```

## #2. Principal Component Analysis

The main PCA function in base R is called `prcomp()`.

Before doing anything like PCA, it is important to check if the data need to be scaled before performing PCA. Recall two common reasons for scaling data include:

-The input variables use different units of measurement. -The input variables have significantly different variances.

```
#checking sd of data
round(apply(wisc.data, 2, sd), 2)
```

radius_mean	texture_mean	perimeter_mean
3.52	4.30	24.30
area_mean	smoothness_mean	compactness_mean
351.91	0.01	0.05
concavity_mean	concave.points_mean	symmetry_mean
0.08	0.04	0.03
fractal_dimension_mean	radius_se	texture_se
0.01	0.28	0.55
perimeter_se	area_se	smoothness_se
2.02	45.49	0.00
compactness_se	concavity_se	concave.points_se
0.02	0.03	0.01
symmetry_se	fractal_dimension_se	radius_worst
0.01	0.00	4.83
texture_worst	perimeter_worst	area_worst
6.15	33.60	569.36
smoothness_worst	compactness_worst	concavity_worst
0.02	0.16	0.21
concave.points_worst	symmetry_worst	fractal_dimension_worst
0.07	0.06	0.02

Looks like we need to scale by setting `scale = TRUE` in our `prcomp()` function call.

#Time for PCA

```
wisc.pr <- prcomp(wisc.data, scale = TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

It's good practice to make a SCREE plot and look for inflection point.

```
attributes(wisc.pr)
```

\$names


```
[1] "sdev"      "rotation" "center"   "scale"    "x"
```

\$class

```
[1] "prcomp"
```


```
pr.var <- wisc.pr$sdev^2
#proportion of variance
pve <- pr.var/sum(pr.var)
plot(pve, xlab = "Principal Component",
```

```
ylab = "Proportion of Variance Explained",  
ylim = c(0, 1), type = "o")
```



Let's make our main results figure from our PCA - our score plot (a.k.a "PC plot", "PC1vsPC2 plot", etc)

```
library(ggplot2)  
pc <- as.data.frame(wisc.pr$x)  
  
ggplot(pc) +  
  aes(PC1, PC2, col = diagnosis) +  
  geom_point()
```



#Hierarchical clustering

Preparation for hierarchical clustering, the distance between all pairs of observations are computed. Furthermore, there are different ways to link clusters together, with single, complete, and average being the most common “linkage methods”.

We can try clustering the original data with `hclut()` or `kmeans()`

First scale `wisc.data` and assign the result to `data.scaled`.

```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)
```

```
head(apply(data.scaled, 2, sd))
```

radius_mean	texture_mean	perimeter_mean	area_mean
1	1	1	1
smoothness_mean	compactness_mean		
1	1		

Calculate the (Euclidean) distance between all pairs of observations in the new scaled dataset and assign the result to `data.dist`.

Create a hierarchical clustering model using complete linkage. Manually specify the method argument to `hclust()` and assign the results to `wisc.hclust`.

```
wisc.hclust <- hclust(dist(data.scaled))  
wisc.hclust
```


Call:

```
hclust(d = dist(data.scaled))
```

```
Cluster method   : complete  
Distance          : euclidean  
Number of objects: 569
```

Let's use the hierarchical clustering model you just created to determine a height (or distance between clusters) where a certain number of clusters exists.

```
plot(wisc.hclust)  
abline(h=19, col="red", lty=2)
```



```
dist(data.scaled)  
hclust (*, "complete")
```

To get a cluster membership vector I will use the `cutree()` function and “cut” into 4 or so grps or clusters.



```
grps <- cutree(wisc.hclust, h=19)
table(grps)
```

```
grps
  1  2  3  4
177  7 383  2
```

I can also use the `table()` to cross tabulate...

```
table(diagnosis)
```

```
diagnosis
  B  M
357 212
```

We can use the `table()` function to compare the cluster membership to the actual diagnoses.

```
table(grps, diagnosis)
```

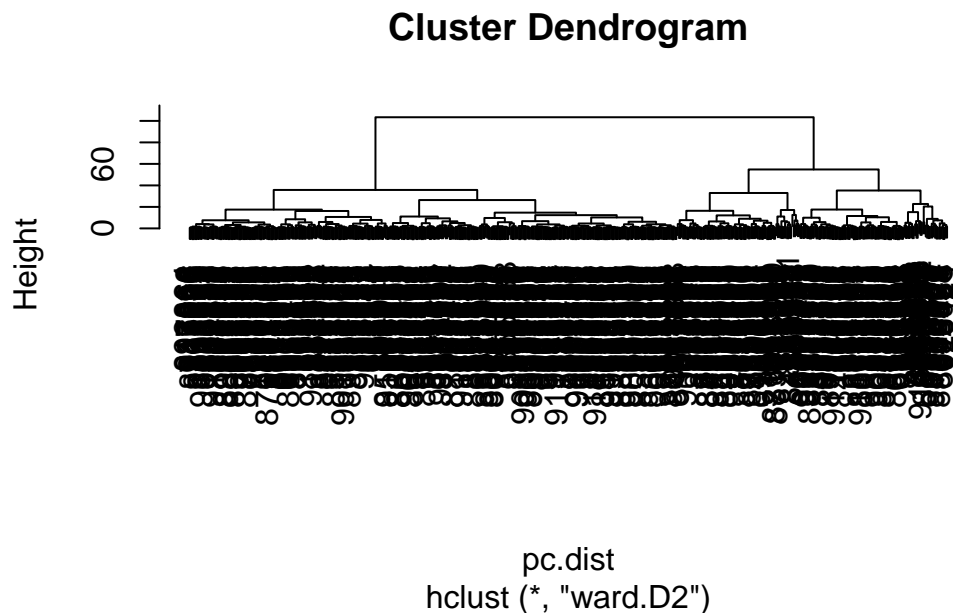
```
      diagnosis
grps  B  M
  1  12 165
  2   2   5
  3 343  40
  4   0   2
```

#Clustering on PCA results

I can cluster in PC-space and use as many or as few PCs as I want.

To start with I will use 3 PCs, that is I will cluster along PC1, PC2, and PC3. Those 3 PCs capture about 70% of variance.

```
#calculate distance and select columns 1:3 for PC1-3.
pc.dist <- dist(wisc.pr$x[,1:3])
#use hclust
wisc.pr.hclust <- hclust(pc.dist, method = "ward.D2")
#plot
plot(wisc.pr.hclust)
```



This looks much nicer than our previous clustering result. Let's find the two major clusters with `cutree` function.

This looks much more promising than our previous clustering results on the original scaled data. Note the two main branches of our dendrogram indicating two main clusters - maybe these are malignant and benign. Let's find out!

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps, diagnosis)
```

```
      diagnosis
grps   B    M
1     24 179
2    333   33
```


#According to our clustering. Cluster 1 groups is associated with malignancy and Cluster 2

We could calculate accuracy - the proportion of samples we got correct if we take cluster 1 to represent all M and cluster 2 to represent all B.


```
(179+333)/nrow(wisc.data)
```

```
[1] 0.8998243
```

```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis)
```



Q.14 How well do the hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.km$cluster` and `wisc.hclust.clusters`) with the vector containing the actual diagnoses

We could calculate accuracy - the proportion of samples we got correct if we take cluster 1 to represent all M and cluster 2 to represent all B.

```
(179+333)/nrow(wisc.data)
```

```
[1] 0.8998243
```

#Sensitivity/Specificity

Sensitivity refers to a test's ability to correctly detect ill patients who do have the condition. In our example here the sensitivity is the total number of samples in the cluster identified as predominantly malignant (cancerous) divided by the total number of known malignant samples. In other words:  $TP/(TP+FN)$ .

#According to our clustering. Cluster 1 groups is associated with malignancy and Cluster 2 groups is associated with benign. When we table our clusters and sort by diagnosis, we notice that in Cluster 1 and Cluster 2 we have patients that are B/M, so false positives or false negatives and this can be used to calculate sensitivity/specificity

```
TP = 179
```

```
FN = 24
```

```
table(grps, diagnosis)
```

```
      diagnosis
grps  B    M
1     24 179
2    333  33
```

```
sensitivity <- TP/(TP+FN)
```

```
sensitivity
```

```
[1] 0.8817734
```

Specificity relates to a test's ability to correctly reject healthy patients without a condition. In our example specificity is the proportion of benign (not cancerous) samples in the cluster identified as predominantly benign that are known to be benign. In other words:  $TN/(TN+FN)$

```
TN = 333
```

```
FN = 33
```

```
specificity <- TN/(TN+FN)
```

```
specificity
```

```
[1] 0.9098361
```

Specificity relates to a test's ability to correctly reject healthy patients without a condition. In our example specificity is the proportion of benign (not cancerous) samples in the cluster identified as predominantly benign that are known to be benign. In other words:  $TN/(TN+FN)$ .

Q15. OPTIONAL: Which of your analysis procedures resulted in a clustering model with the best specificity? How about sensitivity?

#Prediction We will use the predict() function that will take our PCA model from before and new cancer cell data and project that data onto our PCA space.

```
#url <- "new_samples.csv"
```


```
url <- "https://tinyurl.com/new-samples-CSV"
```

```
new <- read.csv(url)
```

```
npc <- predict(wisc.pr, newdata=new)
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q16. Which of these new patients should we prioritize for follow up based on your results? patient 2